



**MEM**  
GRAPH

Paulina Gacek, Jakub Hulek

# Technology Stack

All performance test were conducted on a machine with following specifications:

**Cpu:** Intel Core i5-14600KF

**Memory:** 32.0 GB 6800 MT/s DDR5

**Hard drive:** Lexar SSD NM790 2TB


**System:** WSL2 UBUNTU 22.04, 20GB RAM, 50 GB swap



Database, graph operations, simple spatial calculations



Data preparation, CLI, most spatial calculations



**docker**  
Compose

Project building and data sharing

# How to run

## Build terminal

```
/ads-2024-project-1$ sudo docker compose up --build
```

## CLI terminal

```
/ads-2024-project-1$ sudo docker compose run --rm manager
```

1. Build the project
2. Run the CLI
3. Add your data to directory “/ads-2024-project-1/data”
4. Its ready! Enter any query you like into the CLI

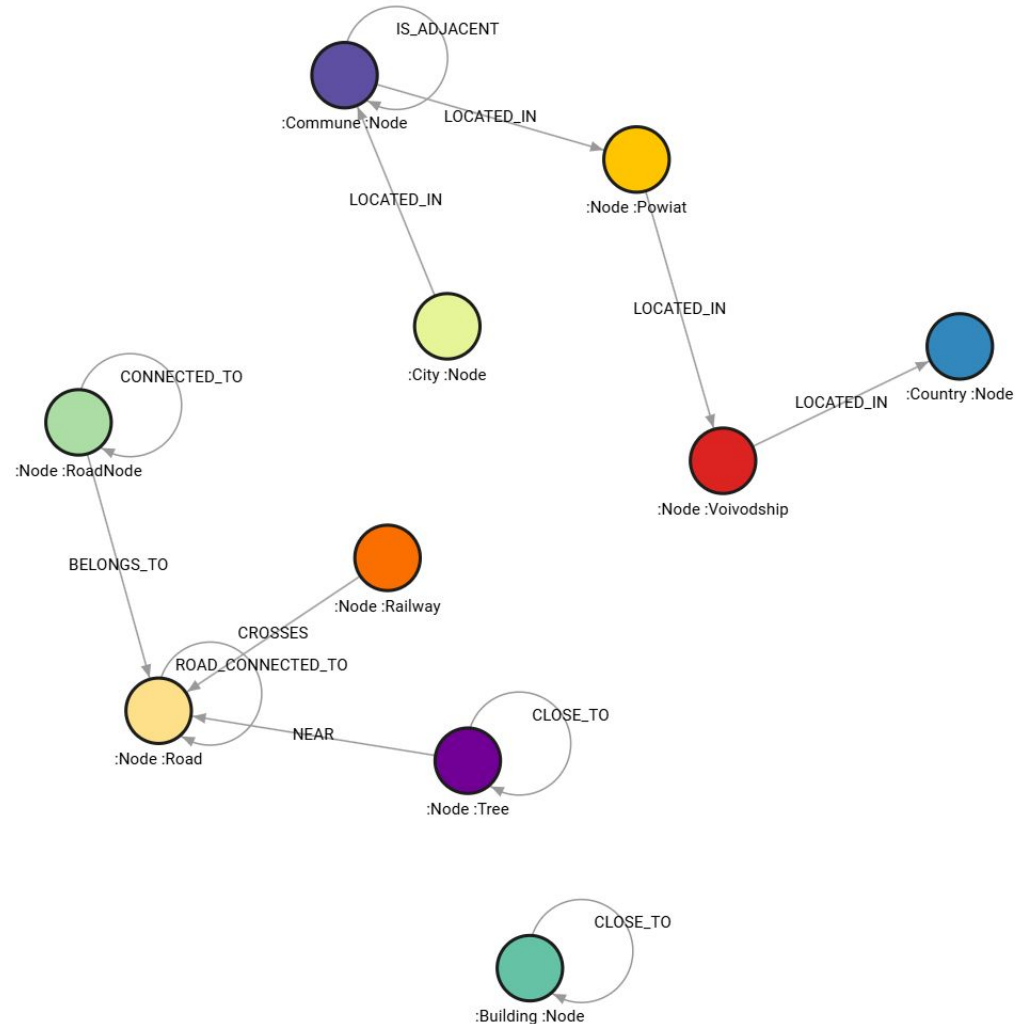
# Graph Schema

All nodes contain float variables “lat” and “lng” for visualisation.

In case of points, spatial data is stored as type “point” in EPSG:2180.

In case of more complicated shapes, spatial data is stored as wkt string in EPSG:2180.

They also include a bounding box as two “point” variables for query optimization.



# Data Import

## Data preparation and data import

Roads data imported in 539.32 seconds.  
Buildings data imported in 693.21 seconds.  
Cities data imported in 18.64 seconds.  
Communes data imported in 7.90 seconds.  
Countries data imported in 0.77 seconds.  
Powiats data imported in 3.76 seconds.  
Railways data imported in 7.90 seconds.  
Trees data imported in 14.39 seconds.  
Voivodships data imported in 1.34 seconds.

**Total time taken: 1287.24 seconds.**  
**Around 21 minutes 37 seconds.**

```
> import auto all > import auto roads trees
```

## Only data import

Roads data imported in 68.20 seconds.  
Buildings data imported in 42.62 seconds.  
Cities data imported in 4.19 seconds.  
Communes data imported in 1.90 seconds.  
Countries data imported in 0.40 seconds.  
Powiats data imported in 1.05 seconds.  
Railways data imported in 2.05 seconds.  
Trees data imported in 3.89 seconds.  
Voivodships data imported in 0.52 seconds.

**Total time taken: 124.82 seconds.**  
**Around 2 minutes 5 seconds**

35.99M

Nodes

46.43M

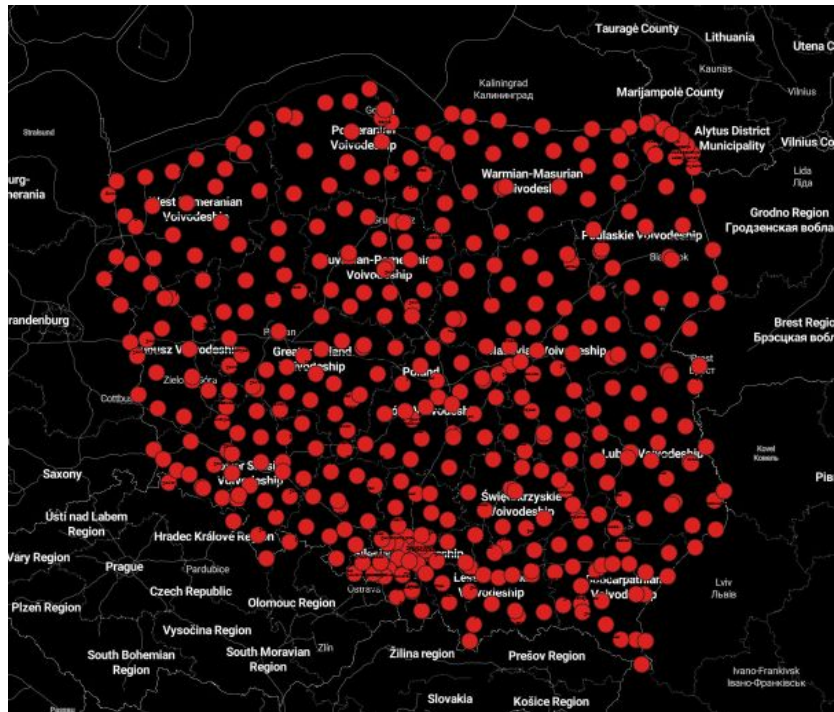
Edges

# Memory usage after Data Import

**23GB**

(25GB peak usage)

### Example road and its nodes



## All powiats in Poland

# Relationship Creation

## Data retrieval, relationship calculation and relationship import

Relationship creation report:

Relationship 1 created in 76.08 seconds.

Relationship 2 created in 9.92 seconds.

Relationship 3 created in 8.82 seconds.

Relationship 4 created in 3.97 seconds.

Relationship 5 created in 46.87 seconds.

Relationship 6 created in 1513.88 seconds.

Relationship 7 created in 27.98 seconds.

Relationship 8 created in 54.51 seconds.

Relationship 9 created in 148.55 seconds.

Relationship 10 created in 167.63 seconds.

**Total time taken: 2058.22 seconds.**

**Around 34 minutes 18 seconds.**

Due to problems with memory usage, we had to switch memory settings to `IN_MEMORY_ANALYTICAL`. This disables ACID compliance, but achieves at least 56B savings on each node and edge and improves performance.

```
> cr all
```

```
> cr 1 2 3 4
```

## Only relationship import (with index creation)

Relationship creation Report:

Relationship 1 created in 41.63 seconds.

Relationship 2 created in 2.00 seconds.

Relationship 3 created in 1.40 seconds.

Relationship 4 created in 1.37 seconds.

Relationship 5 created in 14.53 seconds.

Relationship 6 created in 175.36 seconds.

Relationship 7 created in 37.89 seconds.

Relationship 8 created in 24.87 seconds.

Relationship 9 created in 154.68 seconds.

Relationship 10 created in 36.32 seconds.

**Total time taken: 490.06 seconds.**

**Around 8 minutes 10 seconds.**

35.99M

Nodes

134.55M

Edges



# Memory usage after Relationship Creation

**36GB**

**(55GB peak usage)**

# Relationship Creation

Relationships 1-6, 8, 10:

Cities in communes

Communes in powiats

Powiats in voivodships

Voivodships in countries

Adjacent communes

Neighbouring buildings

Trees close to roads

Railways crossing roads

retrieval optimization with area of interest, calculated with shapely in manager created with "LOAD CSV"

Relationship 7: Neighbouring trees

calculated and created in memgraph, optimized with a POINT INDEX

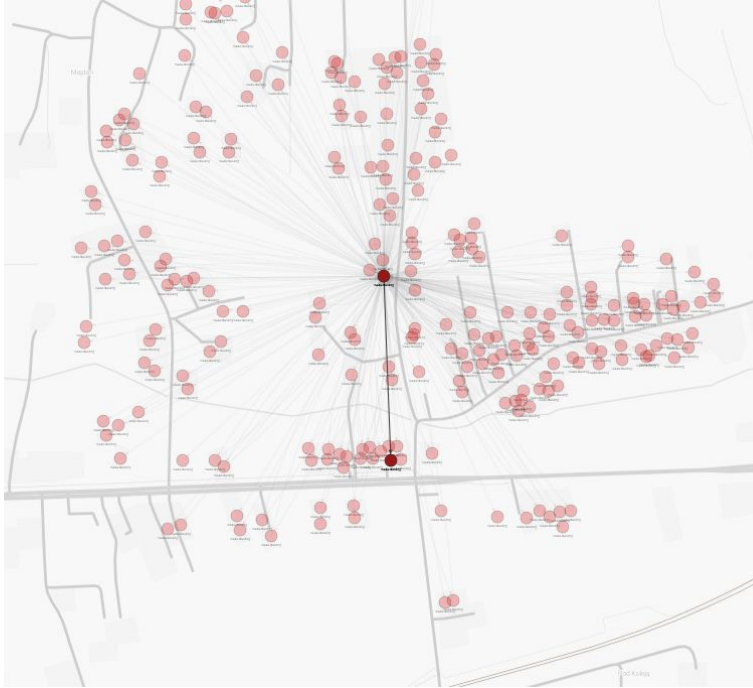
Relationship 9: Connected roads

calculated and created in memgraph




Due to problems with computational complexity neighbouring buildings detection had to be clipped to a small area of Poland. Bounding box of Powiat Wielicki was chosen (140 000 buildings, less than 1% of all buildings). This resulted in *only* 23 million new edges.

# Relationship Creation Examples



Example building neighbours



Example Powiats in Mazowieckie voivodship

# Querying

```
> q all
```

```
> q 1 preset
```

```
> q 4 60 house 10
```

## Creating indexes, running queries

Query running report:

Query 1 run in 2.06 seconds.

Query 2 run in 1.38 seconds.

Query 3 run in 1.03 seconds.

Query 4 run in 102.21 seconds.

Query 5 run in 1.47 seconds.

Query 6 run in 161.85 seconds.

Query 7 run in 90.03 seconds.

Query 8 run in 5.16 seconds.

Query 9 run in 163.84 seconds.

Query 10 run in 5.62 seconds.

**Total time taken: 534.65 seconds.**

**Around 8 minutes 54 seconds.**

## Only running queries

Query running report:

Query 1 run in 0.88 seconds.

Query 2 run in 0.78 seconds.

Query 3 run in 0.72 seconds.

Query 4 run in 60.12 seconds.

Query 5 run in 0.55 seconds.

Query 6 run in 150.70 seconds.

Query 7 run in 33.43 seconds.

Query 8 run in 0.99 seconds.

Query 9 run in 153.68 seconds.

Query 10 run in 3.29 seconds.

**Total time taken: 405.17 seconds.**

**Around 6 minutes 45 seconds**



All queries (apart from query 6 - parallel roads) are executed with a CYPHER query, and results is transformed to an appropriate json form in the manager.

# Example Querying

Query 10 - Trees within 15 meters of a road, at least 10

```
> q 10 10 15
```

```
MATCH (road:Road)
MATCH p=(tree:Tree)-[e:CLOSE_TO]->(road)
WHERE e.distance <= {max_distance}
WITH COLLECT(tree.id) as trees, road
WHERE size(trees) >= {min_count}
RETURN road.id as road_id, road.name as road_name, trees as tree_ids
```




# Example Querying

Query 4 - Clusters of building of type “house”,  
distance 60 meters, at least 10 buildings for a cluster

```
> q 4 60 house 10
```

```
MATCH p=(:Building {{building:"{building_type}"}})-[e:CLOSE_TO]
-(:Building {{building:"{building_type}"}})
WHERE e.distance <= {max_distance}
WITH project(p) AS subgraph
CALL nxalg.strongly_connected_components(subgraph)
YIELD components
UNWIND components as c
WITH c
WHERE size(c) >= {min_count}
RETURN EXTRACT(n in c | n.id) as ids
```





**Thank you!**

Paulina Gacek, Jakub Hulek

